



Floor pricing, with a human in charge

A floor-pricing partner for ad ops. Statistics propose, an agent can refuse, a human decides.

Price floors are set by hand, on a belief, with no way to see the curve

Too low

Money left on the table in every auction

Too high

Auctions go unfilled and revenue drops

No signal

A bidder outage, demand surge or reporting glitch looks like a trend

In sample data, modelled manual pricing captures 78.7% of the available revenue and makes a harmful move on 16.3% of decisions.

Three parts, each doing only what it is good at

1

Statistics layer

Fits revenue per 1k requests against log floor over 21 days. Step is capped and shrinks when confidence is low.

2

Agent that can refuse


Reads ops notes and data-quality signals. Returns a trust plan, never a price. Refuses on thin traffic, flat floors, outlier days.

3


Human decides

Every recommendation sits in a queue with its evidence. Ops accepts, edits or rejects, and the reason is logged.

A working web app and a clickable demo



Floor desk: evidence beside every recommendation



Web app: decision log with reasons, CSV in and out

What exists today, all verifiable in the repo

40 x 180

sample cells and days, with events and ops notes

6

arms evaluated, scored on a held-out world

100


test runs per pilot design, plus A/A placebo

11

automated tests, CI passing on GitHub

Next: a matched-cell pilot with real ops colleagues. Success criteria are already written and checked on sample data.

The statistics layer does the work. Refusal buys safety.



Harmful moves

Ops alone 16.3%. Unguarded stats 6.2%.
Guarded 2.6% to 3.0%.

What did not hold

The agent stand-in did not beat a plain rule. The pair did not beat the model alone. Both stay open until the Claude backend is run.

The findings changed the design four times

Scope

Started as a VAST/GAM discrepancy adjudicator, narrowed to floor pricing where the decision and the outcome are both measurable

v1

Hard significance gate refused 82% of days and cost 1.7 points. Replaced by a confidence-scaled step

Metric

Exploration jitter made held floors look harmful. Harm redefined: move over 10% and revenue down over 1%

Findings

Refusing thin cells freezes learning. Holding during an outage can hurt

Pilot

Tested the pilot design itself: 20 cells for one week passes 26% of the time. Redesign to 40 cells, 4 live weeks

One builder, product and engineering in one head

J John Britto

Head of Product, Bengaluru

- 8+ years in AdTech, from UI/UX designer to Head of Product
- Built ad servers, header bidding, OpenRTB exchange and reporting pipelines
- Leads backend, frontend and data engineers, and the roadmap
- PLUMB built solo: problem, statistics, agent design, evaluation, app



Statistics propose. The agent can refuse. A human decides.

Demo: plumb-floor-pricing-demo.vercel.app

Code: github.com/Johnbritt/Plumb-Floor-Pricing